

# Practical Reinforcement Learning

Episode 3.5

## Deep Learning 101



Yandex  
Data Factory


LAMBDA



British Hedgehog  
Preservation Society



# Recap: logistic regression



# Recap: Gradient descent


$$P(y|x) = \sigma(w \cdot x + b)$$

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

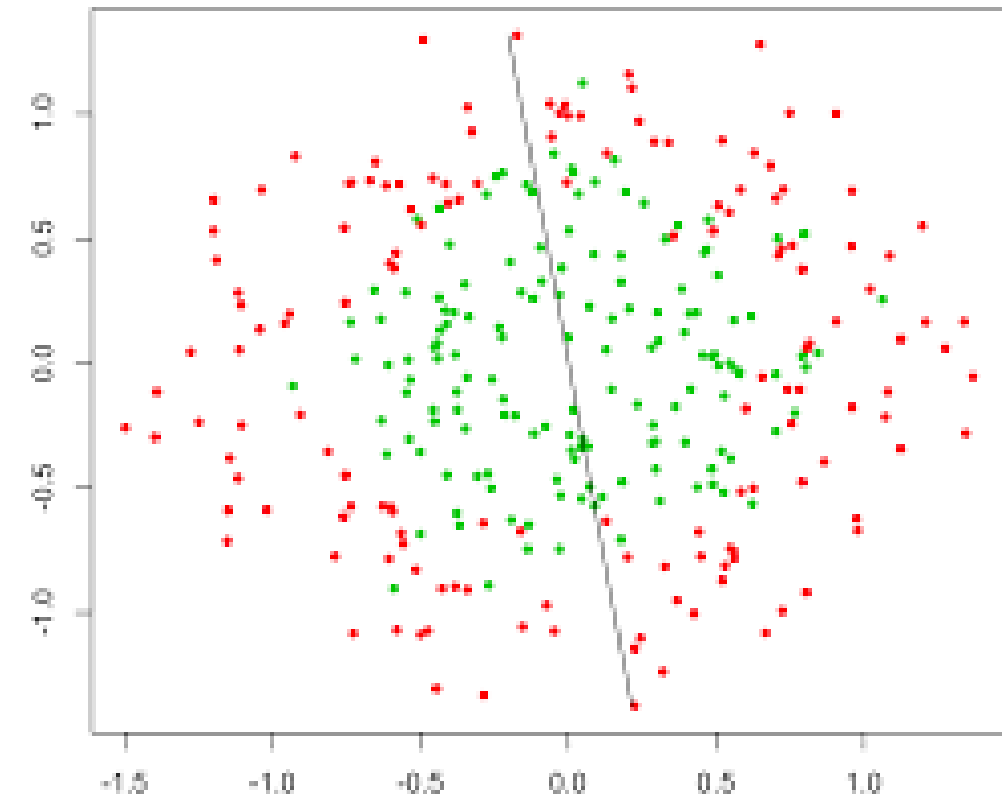
Repeat until convergence

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial L(y, y_{pred})}{\partial \theta_j}$$


$$\Theta \sim \{w, b\}$$



# Problem: Nonlinear dependencies



What we have



What we want


- **Trivia:** how do we solve that with logistic regression?

# Feature extraction

Loss, for example:


$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Model:



Training:

$$\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$$




Features would tune to your problem automatically!

# What do we want, exactly?

Loss, for example:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Model:




Training:                      ?                       $\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$

# What do we want, exactly?

Loss, for example:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Model:




Gradients:  $\underset{\theta_2}{\operatorname{argmin}} L(y, P(y|x))$      $\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$



# Try linear

Model:




$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

# Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:

$$P(y|x) = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

Is it any better than logistic regression?

# Try linear


$$P(y|x) = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

$$w'_i = \sum_j w_j^o w_{ij}^h \qquad b' = \sum_j w_j^o b_j^h + b^o$$

$$P(y|x) = \sigma\left(\sum_i w'_i x_i + b'\right)$$

# Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$


Output:

$$P(y|x) = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

Is it any better than logistic regression?

# Nonlinearity

Model:

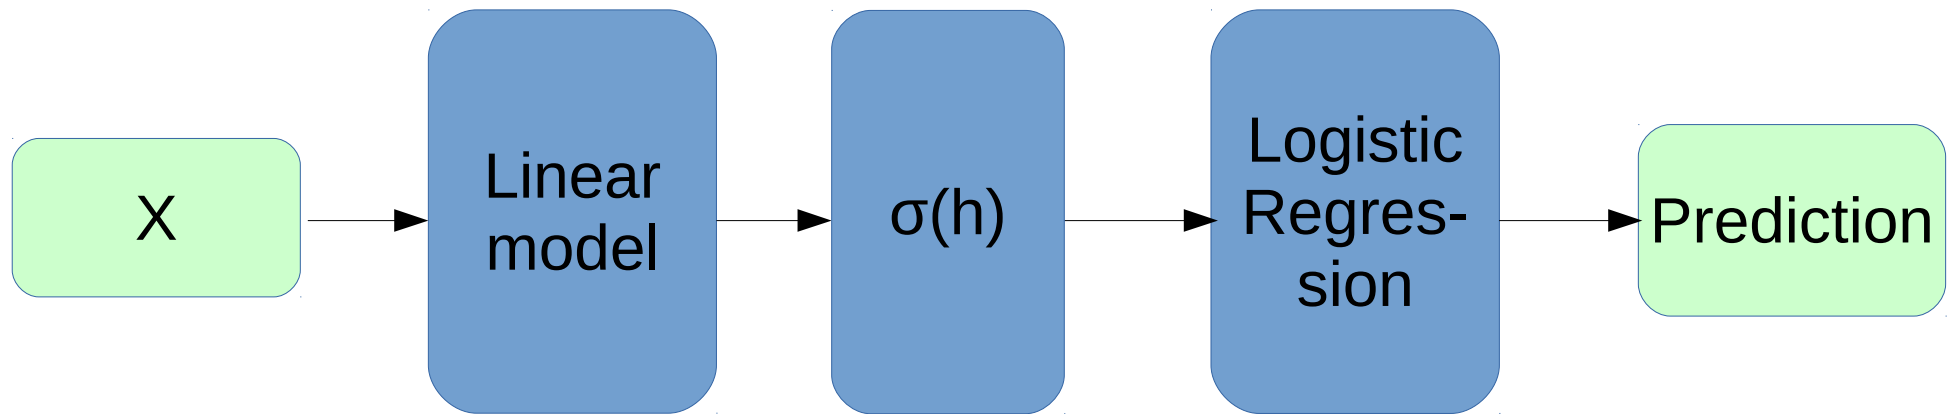


$$h_j = \sigma\left(\sum_{j \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

# Nonlinearity

Model:



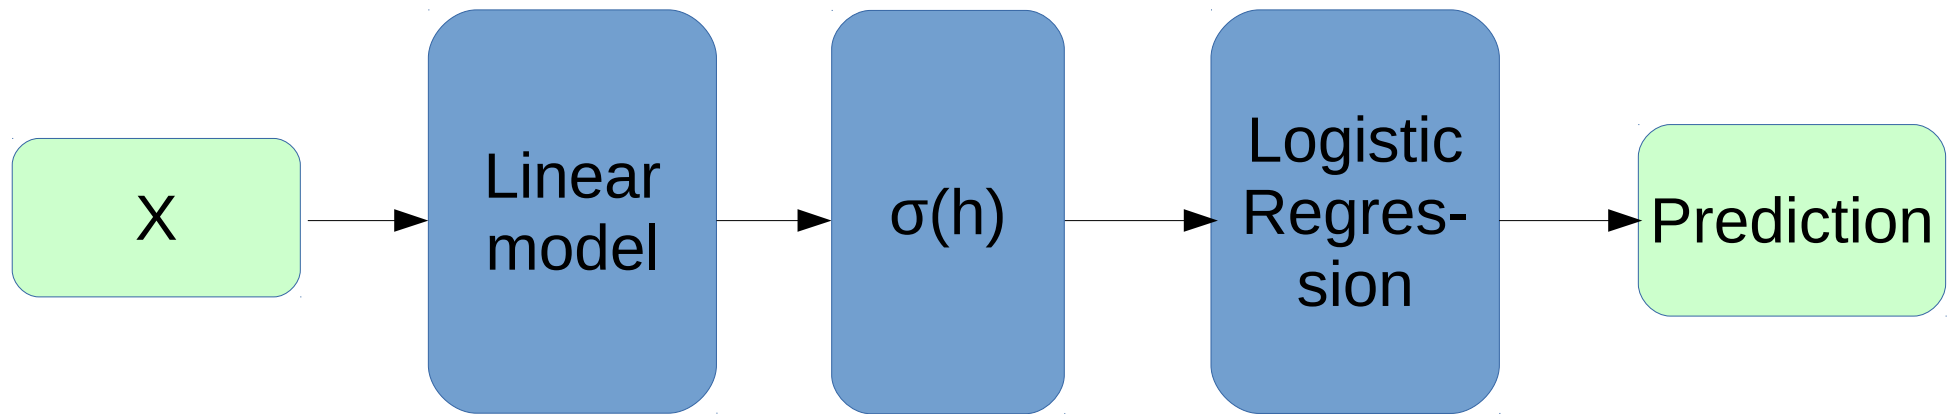
$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right) \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:

$$P(y|x) = \sigma\left(\sum_j w_j^o \sigma\left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$


# Nonlinearity

Model:



$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$


$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$



# Nonlinearity

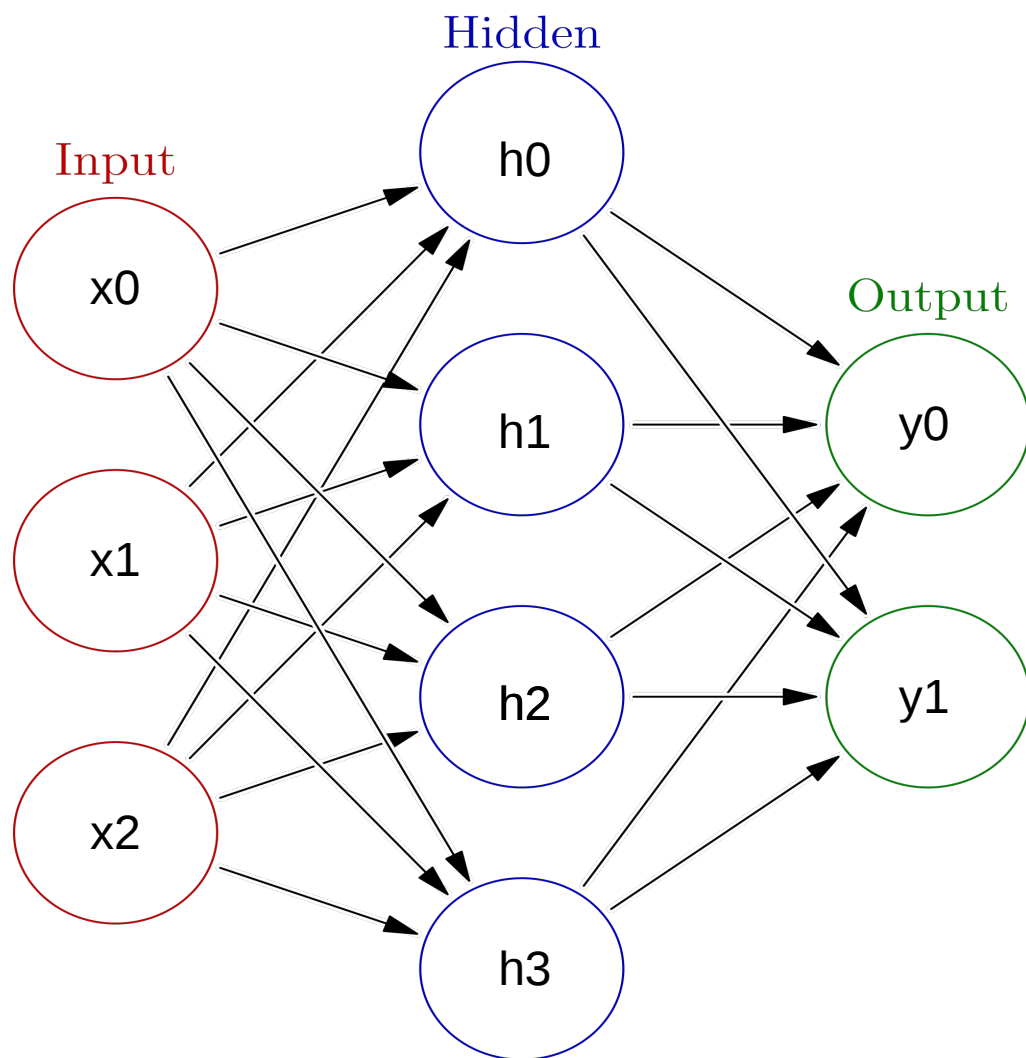
- $f(a) = 1/(1+e^a)$
- $f(a) = \tanh(a)$

- $f(a) = \max(0, a)$
- $f(a) = \log(1+e^a)$






# Initialization, symmetry problem



- Initialize with zeros  
 $W \leftarrow 0$
- What will the first step look like?

# Initialization, symmetry problem




- Break the symmetry!
- Initialize with random numbers!  
 $W \leftarrow N(0, 0.01)?$   
 $W \leftarrow U(0, 0.1)?$
- Can get a bit better for deep NNs

# Connectionist phrasebook


- Layer – a building block for NNs :
  - “Dense layer”:  $f(x) = Wx + b$
  - “Nonlinearity layer”:  $f(x) = \sigma(x)$
  - Input layer, output layer
  - A few more we gonna cover later
- Activation – layer output
  - i.e. some intermediate signal in the NN
- Backpropagation – a fancy word for “chain rule”

# Connectionist phrasebook



- “Train it via backprop!”

# Connectionist phrasebook



How do we train it?

# Image recognition



“Dog”

# Image recognition



“Gray wall”

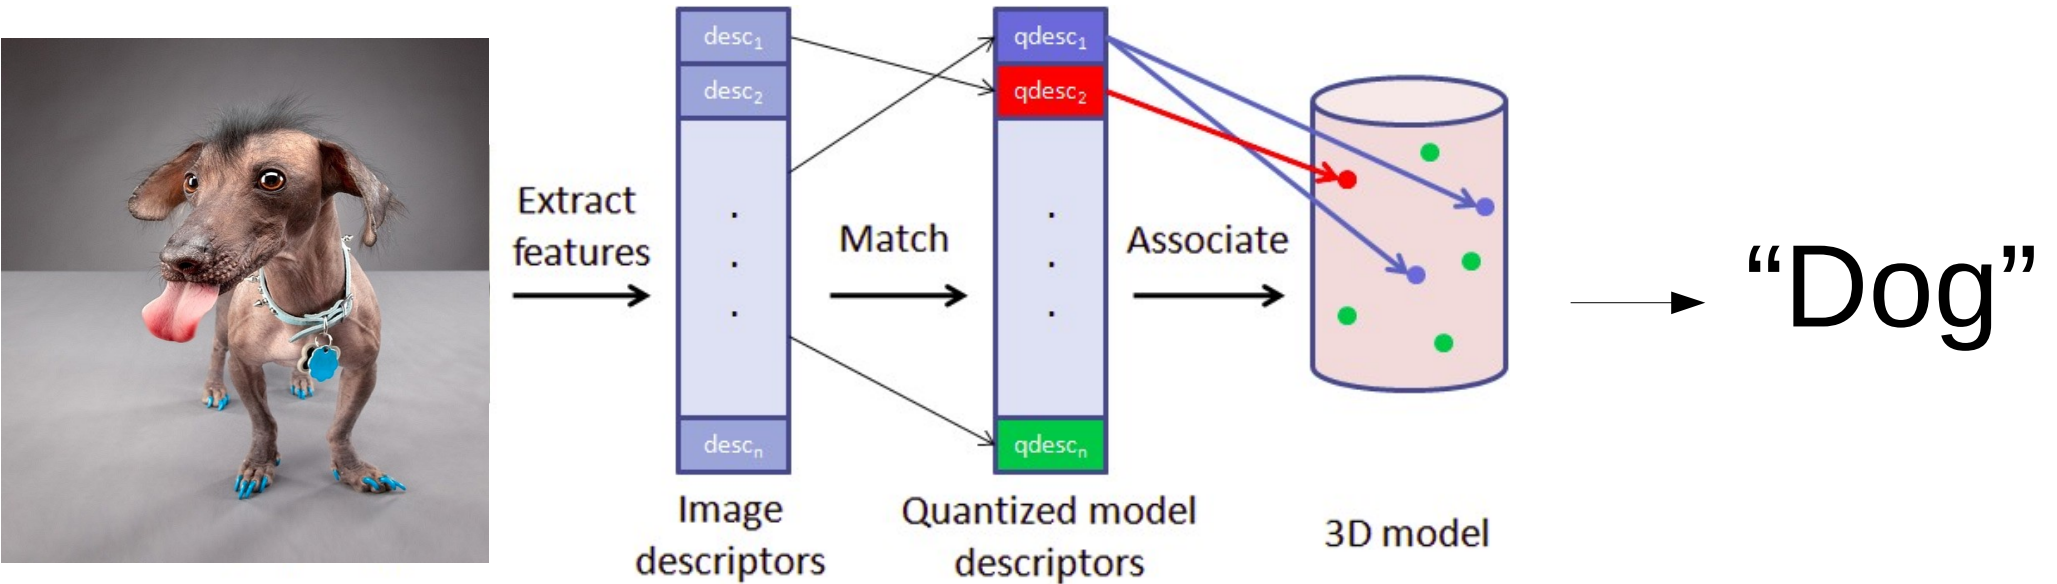
“Dog tongue”

“Dog”

<a particular kind  
of dog>


“Animal sadism”

# Classical approach





# NN approach




**What features could NN learn this way?**




**Discrete Choices**


⋮



**Layer 2 Features**



**Layer 1 Features**



**Original Data**

# Problem


Should we require, say, “Dog ear” feature

- Linear combination can only select dog ear at a one (or a few) positions.
- Need to learn independent features for each position
- Next layer needs to react on “dog ear 0,0 or dog ear 0,1 or ... or dog ear 255,255”
- Introduce **a lot** of parameters and risk overfitting.


Idea: force all these “dog ear” features to use **exactly same weights**, shifting weight matrix each time.

# Convolution

- Apply same weights to all patches



# Convolution



apply same filter to all patches

# Convolution

5x5

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image


3x3 (5-3+1)

4		

Convolved  
Feature

Intuition: how cat-like is this square?


# Convolution



Intuition: how cat-like is this square?

# Convolution

Input image



Convolution  
Kernel

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Feature map




Intuition: how **edge-like** is this square?



# Convolution



Image : 3 (RGB) x 100 px x 100 px




Filters: 100x(3x5x5)

# Convolution



Image : 3 (RGB) x 100 px x 100 px




Filters: 100x(3x5x5)

100x96x96  
~10<sup>6</sup>

Somewhat too many!

# Pooling



Intuition: What is the max cat-likelihood over this area?


# Pooling

Motivation:

- Reduce layer size by a factor
- Make NN less sensitive to small image shifts

Popular types:


- Max
- Mean(average)






**Discrete Choices**


⋮



**Layer 2 Features**




**Layer 1 Features**




**Original Data**

# Convolutional NNs




# Convolutional NNs



## Quiz:

1) What is the blob size **after second pooling**

# Convolutional NNs




## Quiz:

2) How many image pixels does **one cell** after **second convolution** depend on?





# Convolutional NNs



## Quiz:

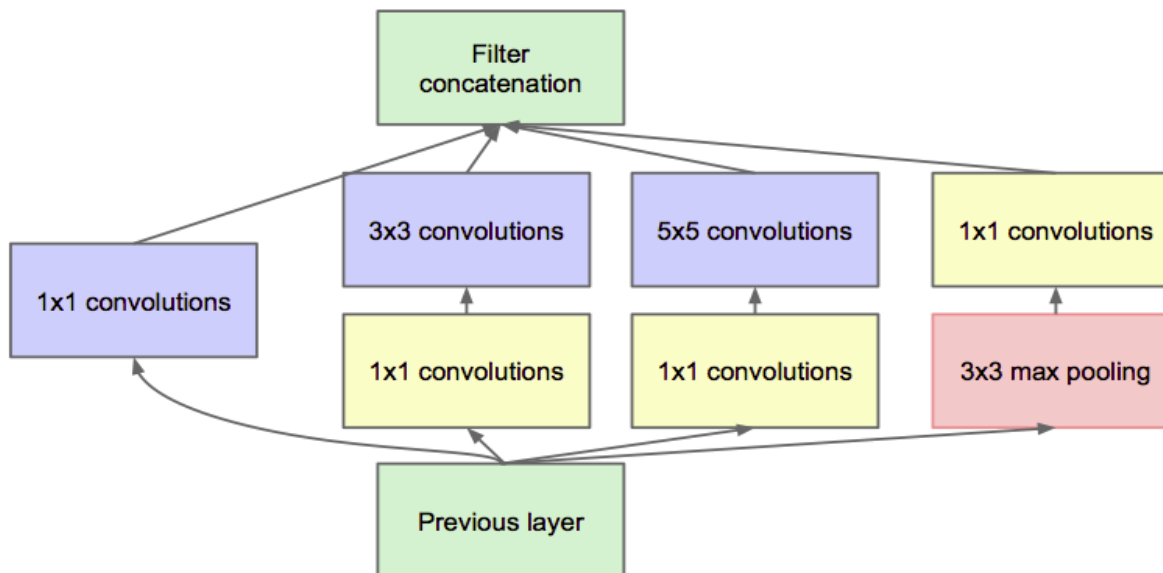

- 3) Which layer is hardest to compute?
- 4) Which layer has most independent parameters?

# Inception-GoogLeNet



It is not a moon. It is a space station (c)

# Inception-GoogLeNet




**Convolution**  
**Pooling**  
**Softmax**  
**Other**

# Potential caveats?


- Hardcore overfitting
- No “golden standard” for architecture
- Computationally heavy

# Regularization

- L1, L2, as usual
- Dropout



(a) Standard Neural Net



(b) After applying dropout.

# Data augmentation



- Idea: we can get N times more data by tweaking images.
- If you rotate cat image by  $15^\circ$ , it's still a cat
- Rotate, crop, zoom, flip horizontally, add noise, etc.
- Sound data: add background noise

# Computation




*The Difference between a CPU and GPU*



**CPU**

MULTIPLE CORES



**GPU**

THOUSAND OF CORES

# Batch normalization

## Problem:

- Consider a neuron in any layer beyond first
- At each iteration we tune it's weights towards better loss function
- But we also tune it's inputs. Some of them become larger, some – smaller
- Now the neuron needs to be re-tuned for it's new inputs



# Batch normalization

TL;DR:

- It's usually a good idea to normalize linear model inputs

(c) Every machine learning lecturer, ever

# Batch normalization

Idea:

- We normalize activation of a hidden layer  
(zero mean unit variance)

$$h_i = \frac{h_i - \mu_i}{\sqrt{\sigma_i^2}}$$

- Update  $\mu_i, \sigma_i^2$  with moving average while training

$$\mu_i := \alpha \cdot \text{mean}_{batch} + (1 - \alpha) \cdot \mu_i$$

$$\sigma_i^2 := \alpha \cdot \text{variance}_{batch} + (1 - \alpha) \cdot \sigma_i^2$$

# Batch normalization

Idea:

- We normalize activation of a hidden layer  
(zero mean unit variance)

$$h_i = \frac{h_i - \mu_i}{\sqrt{\sigma_i^2}}$$


i stands for i-th neuron

- Update  $\mu_i, \sigma_i^2$  with moving average while training


$$\mu_i := \alpha \cdot \text{mean}_{\text{batch}} + (1 - \alpha) \cdot \mu_i$$

$$\sigma_i^2 := \alpha \cdot \text{variance}_{\text{batch}} + (1 - \alpha) \cdot \sigma_i^2$$

# Batch normalization




VS



# Batch normalization

## Good side effect #1:


- Vanishing gradient less a problem for sigmoid-like nonlinearities



## Good side effect #2:

- We no longer need to train bias (+b term in  $Wx+b$ )

# Batch Normalization



# More

## **Regularization:**

- dropconnect, variational dropout, ...

## **Normalization:**

- weight normalization,
- normalization propagation,
- layer normalization,

## **Initialization:**

- data-aware initialization, pre-training, ...

# Nuff

**Let's code some neural networks!**

