

Reinforcement learning

Episode 6

Policy gradient methods



Yandex
Data Factory

LAMBDA 



**British Hedgehog
Preservation Society**

Small experiment

The next slide contains a question

Please respond as fast as you can!

Small experiment



left or right?

Small experiment




Right! Ready for next one?

Small experiment



What's $Q(s, \text{right})$ under $\gamma=0.99$?

Small experiment



What's $Q(s, \text{right})$ under $\gamma=0.99$?

Conclusion

- Often computing q-values is harder than picking optimal actions!
- We could avoid learning value functions by directly learning agent's policy $\pi_{\theta}(a|s)$

Trivia: what algorithm works that way?

(of those we studied)

Conclusion

- Often computing q-values is harder than picking optimal actions!
- We could avoid learning value functions by directly learning agent's policy $\pi_{\theta}(a|s)$

Trivia: what algorithm works that way?

e.g. crossentropy method

NOT how humans survived

$\text{argmax}[$
Q(s, pet the tiger)
Q(s, run from tiger)
Q(s, provoke tiger)
Q(s, ignore tiger)
]



how humans survived

$$\pi(run|s)=1$$



Policies

In general, two kinds

- Deterministic policy

$$a = \pi_{\theta}(s)$$

- Stochastic policy

$$a \sim \pi_{\theta}(a|s)$$

Trivia: Any reason why stochastic is better?

Policies

In general, two kinds

- Deterministic policy

$$a = \pi_{\theta}(s)$$

- Stochastic policy

$$a \sim \pi_{\theta}(a|s)$$

Trivia: Any reason why stochastic is better?

Policies

In general, two kinds

- Deterministic policy

$$a = \pi_{\theta}(s)$$

- Stochastic policy

$$a \sim \pi_{\theta}(a|s)$$

Trivia: Any reason why stochastic can be better?


Policies

In general, two kinds

- Deterministic policy

Genetic algos (week 0)
Deterministic policy gradient


same action each time


$$a = \pi_{\theta}(s)$$

- Stochastic policy

Crossentropy method
Policy gradient

sampling takes care
of exploration


$$a \sim \pi_{\theta}(a|s)$$


Policies

In general, two kinds

- Deterministic policy

Genetic algos (week 0)
Deterministic policy gradient


same action each time


$$a = \pi_{\theta}(s)$$

- Stochastic policy

Crossentropy method
Policy gradient

sampling takes care
of exploration


$$a \sim \pi_{\theta}(a|s)$$

categorical, normal, mixture of normal, whatever

Two approaches

- **Value based:**

Learn value function $Q_{\theta}(s, a)$ or $V_{\theta}(s)$

Infer policy $\pi(a|s) = [a = \underset{a}{\operatorname{argmax}} Q_{\theta}(s, a)]$

- **Policy based:**

Explicitly learn policy $\pi_{\theta}(a|s)$ or $\pi_{\theta}(s) \rightarrow a$

Implicitly maximize reward over policy

Recap: crossentropy method

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions
 - elite = take M best sessions and concatenate

$$\theta_{i+1} = \theta_i + \alpha \nabla \sum_i \log \pi_{\theta_i}(a_i | s_i) \cdot [s_i, a_i \in \text{Elite}]$$

Trivia: Can we adapt it for discounted rewards?
(with γ)

Recap: crossentropy method

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions
 - elite = take M best sessions and concatenate

$$\theta_{i+1} = \theta_i + \alpha \nabla \sum_i \log \pi_{\theta_i}(a_i | s_i) \cdot [s_i, a_i \in \text{Elite}]$$

TD version: elite (s,a) that have highest R(s,a)
(select elites independently from each state)

Policy gradient main idea

Why so complicated?

We'd rather simply maximize R over π !

Objective

Expected reward:

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a) \\ \dots}} R(s, a, s', a', \dots)$$

Expected discounted reward:

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} Q(s, a)$$

Objective

Expected reward: $R(z)$ setting

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a) \\ \dots}} R(s, a, s', a', \dots)$$

Expected discounted reward: $R(s,a) = r + \gamma * R(s',a')$

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} Q(s, a)$$

“true” Q-function



Objective

$$J = \mathop{E}_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} Q(s, a) = \int_s p(s) \int_a \pi_\theta(a|s) Q(s, a) da ds$$

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} Q(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) Q(s, a) da ds$$

Diagram illustrating the components of the objective function J :

- Agent's policy**: Points to $\pi_{\theta}(a|s)$.
- state visitation frequency (may depend on policy)**: Points to $p(s)$.
- True action value**: Points to $Q(s, a)$.

Trivia: how do we compute that?

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} Q(s, a) = \int_s p(s) \int_a \pi_\theta(a|s) Q(s, a) da ds$$

$$J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in z_i} Q(s, a)$$

Empirical action value
a.k.a. $R(s, a)$

sample N sessions

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} Q(s, a) = \int_s p(s) \int_a \pi_\theta(a|s) Q(s, a) da ds$$

$$J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in z_i} Q(s, a)$$

Empirical action value
a.k.a. $R(s, a)$

sample N sessions

Can we optimize policy now?

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} Q(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) Q(s, a) da ds$$

parameters “sit” here

$$J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in z_i} Q(s, a)$$

Empirical action value
a.k.a. $R(s, a)$

We don't know how to compute $dJ/d\theta$

Optimization

Finite differences

- Change policy a little, evaluate

$$\nabla J \approx \frac{J_{\theta+\epsilon} - J_{\theta}}{\epsilon}$$

Stochastic optimization

- Good old crossentropy method
- Maximize probability of “elite” actions

Optimization

Finite differences

- Change policy a little, evaluate

$$\nabla J \approx \frac{J_{\theta+\epsilon} - J_{\theta}}{\epsilon}$$

Stochastic optimization

- Good old crossentropy method
- Maximize probability of “elite” actions

Trivia: any problems with those two?

Optimization

Finite differences

- Change policy a little, evaluate

$$\nabla J \approx \frac{J_{\theta+\epsilon} - J_{\theta}}{\epsilon}$$

VERY noisy, especially
if both J are sampled

Stochastic optimization

- Good old crossentropy method
- Maximize probability of “elite” actions

“quantile convergence”
problems with stochastic
MDPs

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} Q(s, a) = \int_s p(s) \int_a \pi_\theta(a|s) Q(s, a) da ds$$

Wish list:

- Analytical derivative
- Easy/stable approximations

Logderivative trick

Simple math

$$\nabla \log \pi(z) = ? ? ?$$

(try chain rule)

Logderivative trick

Simple math


$$\nabla \log \pi(z) = \frac{1}{\pi(z)} \cdot \nabla \pi(z)$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$

Policy gradient

Analytical inference


$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) Q(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


Policy gradient

Analytical inference

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) Q(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


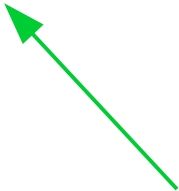
$$\nabla J = \int_s p(s) \int_a \pi_\theta(a|s) \nabla \log \pi_\theta(a|s) Q(s, a) da ds$$

Trivia: anything curious about that formula?

Policy gradient

Analytical inference

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) Q(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


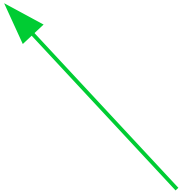
$$\nabla J = \int_s p(s) \int_a \pi_\theta(a|s) \nabla \log \pi_\theta(a|s) Q(s, a) da ds$$

that's expectation :)

Policy gradient

Analytical inference

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) Q(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


$$\nabla J = E_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

Policy gradient (REINFORCE)

- Policy gradient

$$\nabla J = E_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

- Approximate with sampling

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in z_i} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

REINFORCE algorithm

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

- Ascend $\theta_{i+1} \leftarrow \theta_i + \alpha \cdot \nabla J$

REINFORCE algorithm

- Initialize NN weights $\theta_0 \leftarrow \text{random}$

Trivia: is it off- or on-policy?

- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

- Ascend $\theta_{i+1} \leftarrow \theta_i + \alpha \cdot \nabla J$

REINFORCE algorithm

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
actions under current policy
= on-policy
 - Evaluate policy gradient

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

- Ascend $\theta_{i+1} \leftarrow \theta_i + \alpha \cdot \nabla J$

value-based Vs policy-based

Value-based

- Q-learning, SARSA, MCTS value-iteration
- Solves harder problem
- Artificial exploration
- Learns from partial experience (temporal difference)
- Evaluates strategy for free :)

Policy-based

- REINFORCE, CEM
- Solves easier problem
- Innate exploration
- Innate stochasticity
- Support continuous action space
- Learns from full session only



value-based Vs policy-based

Value-based

- Q-learning, SARSA, MCTS value-iteration
- Solves harder problem
- Artificial exploration
- Learns from partial experience (temporal difference)
- Evaluates strategy for free :)

Policy-based

- REINFORCE, CEM
- **We'll learn much more soon!**
- Solves easier problem
- Innate exploration
- Innate stochasticity
- Support continuous action space
- ~~Learns from full session only~~



REINFORCE algorithm

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

What is better for learning:
random action in good state
or
great action in bad state?

REINFORCE baseline

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient


$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot Q(s, a)$$

$$Q(s, a) = V(s) + A(s, a)$$

Actions influence $A(s, a)$ only, so $V(s)$ is irrelevant

REINFORCE baseline

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot (Q(s, a) - b(s))$$


Anything that doesn't depend on action
ideally, $b(s) = V(s)$

Actor-critic

- Learn both $V(s)$ and $\pi_{\theta}(a|s)$
- Hope for best of both worlds :)



Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$

Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

Non-trivia: how can we estimate $\mathbf{A(s,a)}$
from (s,a,r,s') and V-function?

Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$

Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

$$A(s, a) = Q(s, a) - V(s)$$

$$Q(s, a) = r + \gamma \cdot V(s')$$

$$A(s, a) = r + \gamma \cdot V(s') - V(s)$$

Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$

Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

$$A(s, a) = Q(s, a) - V(s)$$

$$Q(s, a) = r + \gamma \cdot V(s')$$

$$A(s, a) = r + \gamma \cdot V(s') - V(s)$$

Also: n-step
version

Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$


Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

$$A(s, a) = r + \gamma \cdot V(s') - V(s)$$

$$\nabla J_{actor} \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in z_i} \nabla \log \pi_{\theta}(a|s) \cdot \underbrace{A(s, a)}_{\text{consider const}}$$

Trivia: how do we train V then?

Advantage actor-critic



Improve policy:

$$\nabla J_{actor} \approx \frac{1}{N} \sum_{i=0}^N \sum_{s,a \in z_i} \nabla \log \pi_{\theta}(a|s) \cdot A(s, a)$$

Improve value:

$$L_{critic} \approx \frac{1}{N} \sum_{i=0}^N \sum_{s,a \in z_i} (V_{\theta}(s) - [r + \gamma \cdot V(s')])^2$$

Continuous action spaces

What if there's continuously many actions?

- Robot control: control motor voltage
- Trading: assign money to equity

How does the algorithm change?

Continuous action spaces

What if there's continuously many actions?

- Robot control: control motor voltage
- Trading: assign money to equity

How does the algorithm change?

it doesn't :)
Just plug in a different formula for
 $\pi(a|s)$, e.g. normal distribution

Duct tape zone

- $V(s)$ errors less important than in Q-learning
 - actor still learns even if critic is random, just slower
- Regularize with entropy
 - to prevent premature convergence
- Learn on parallel sessions
 - Or super-small experience replay
- Use logsoftmax for numerical stability



Let's code!